



# CNN Applications

CS5242

Lee Hwee Kuan & **Wang Wei**

Teaching assistant:

Connie Kou Khor Li, Ji Xin, Ouyang Kun

[cs5242@comp.nus.edu.sg](mailto:cs5242@comp.nus.edu.sg)

# Recap

- Operations
  - Convolution
  - Pooling
  - Dropout
  - ReLU
  - Batch-Normalization
  - Fully Connected
  - Softmax
  - Cross-entropy

# Recap

- Architectures for ImageNet classification
  - AlexNet
  - VGG
  - InceptionNet
  - ResNet

# Roadmap

## CNN applications

- Image classification
- Object detection
- Object segmentation

## Attention modelling

# Image classification

- Predict the class/label of the image
- Training label
  - ground truth label (index)
- Test output
  - A probability distribution vector, one probability per label

*Source from [13]*



Training label: **bicycle**  
Prediction output:  
**bicycle 0.6; people 0.3; mountain 0.05;**

# Image classification

- Approaches
  - AlexNet, VGG, InceptionNet, ResNet, DenseNet, etc
  - With a Softmax layer as the final output layer
  - With cross-entropy as the loss function
- Dataset
  - ImageNet
- Evaluation
  - Top-1: accuracy =  $\#(\text{top1 prediction is truth label}) / \# \text{ test samples}$
  - Top-5: accuracy =  $\#(\text{one of top5 prediction is truth label}) / \# \text{ test samples}$

# Applications

- [Logo classification](#)
- [Traffic sign classification](#)
  - [notebook](#)
- [Ecommerce product classification](#)
- Medical image classification
- Food image classification
- ImageNet classification
- Dogs vs Cats
  - [notebook](#)

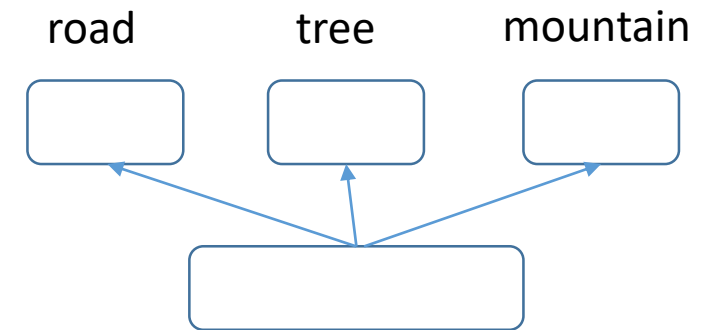
# Image annotation

- Approaches

- Same architecture as image classification
- Binary classification for each label (0 vs 1)
  - Logistic function as the output layer
  - Cross-entropy for each label
- Other output and loss layers [1]

- Evaluation [1]

- Precision = average over all test samples  $\{ | \text{Prediction} \cap \text{Truth} | / \# \text{Prediction} \}$
- Recall = average over all test samples  $\{ | \text{Prediction} \cap \text{Truth} | / \# \text{Truth} \}$





# Image annotation

- Application
  - Example: [satellite image annotation](#)
  - [Notebook](#)



Source from: <https://www.kaggle.com/c/planet-understanding-the-amazon-from-space>

# Object detection

- Detect the location of all object instances of all classes
- Training label
  - a list of <class, bounding box of each object instance>
- Prediction output
  - a list of <class, probability, bounding box of each object instance>



Training label:

bicycle (10, 100, 110, 110) (200, 200, 180, 80) ...

People (200, 80, 71, 71) (300, 50, 20, 80) ...

Prediction output:

bicycle 0.9 (9, 93, 100, 111), 0.8 (200, 200, 180, 80), ...

people 0.8 (200, 80, 71, 71), ...

# Object detection

- Applications
  - Face detection
    - Point-and-shoot camera
  - Surveillance
    - Count cars, peoples, animals
  - Indexing
    - Get objects from images for search
- Evaluation
  - Matched prediction = detected bounding box has enough overlap with truth and its label is correct
  - Precision =  $\text{\#matched prediction} / \text{\#total predictions}$
  - Recall =  $\text{\#matched prediction} / \text{\#truth instance (bounding box)}$

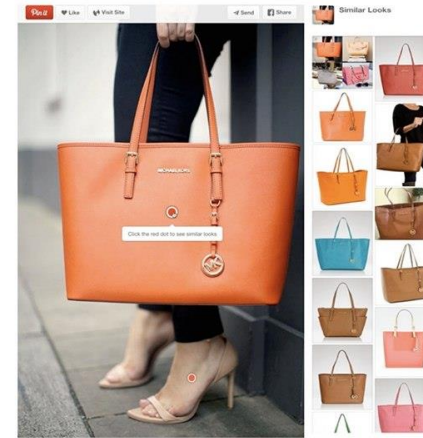


Figure 1: Similar Looks: We apply object detection to localize products such as bags and shoes. In this prototype, users click on automatically tagged objects to view similar-looking products.

Source from [10]

# Object detection

- Solution
  - Find some candidate regions with objects
  - Extract CNN feature from this region
  - Refine the region boundary (bounding box) using a regressor
    - Generate 4 values (x, y, h, w)
  - Predict the class label using Softmax

# R-CNN [7]

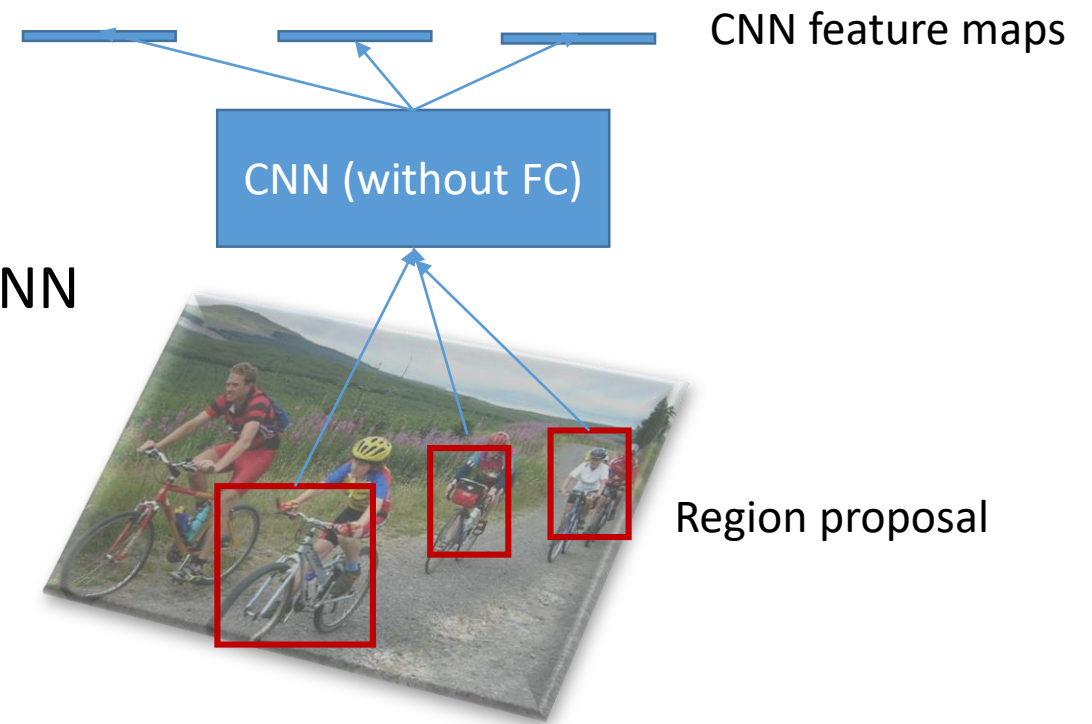
- Generate region proposal
  - Candidate object regions
  - Using existing methods



Region proposal

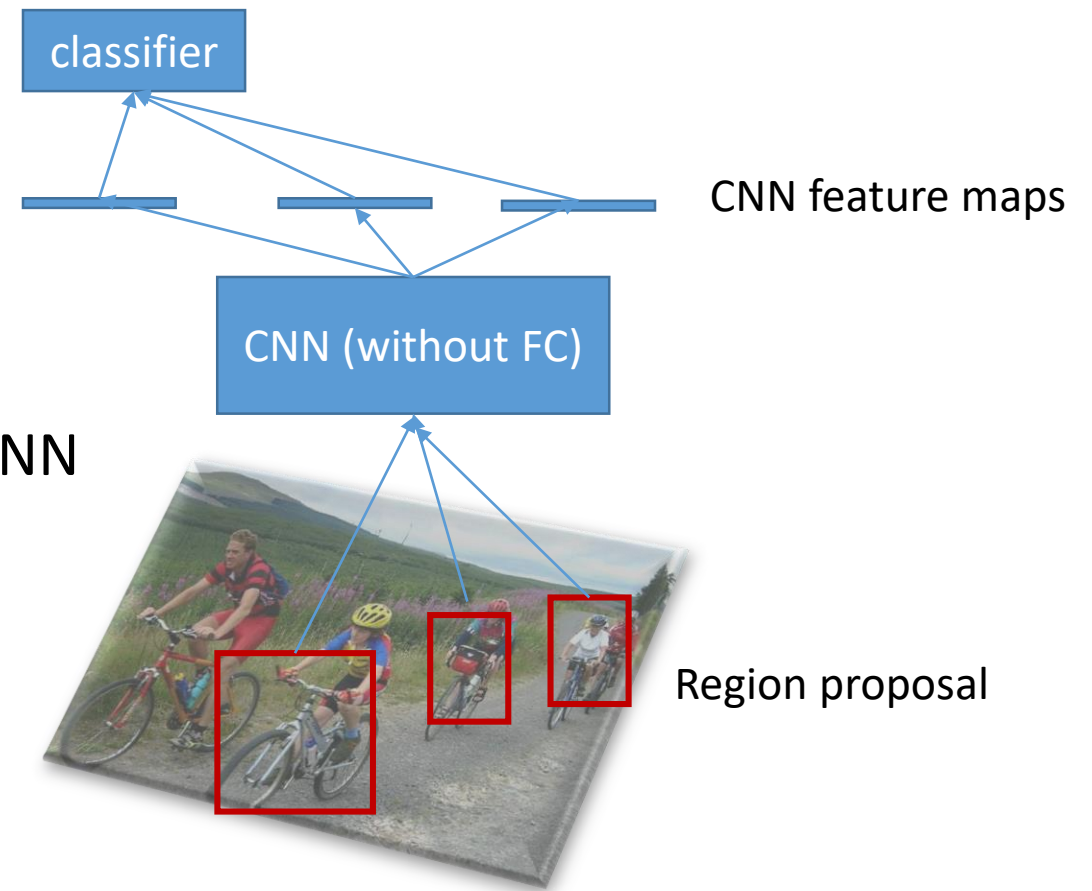
# R-CNN

- Generate region proposal
  - Candidate object regions
  - Using existing methods
- Extract CNN feature
  - For each region
  - Forward-propagate each region via CNN
  - Using popular CNN architecture
    - VGG/ResNet/etc



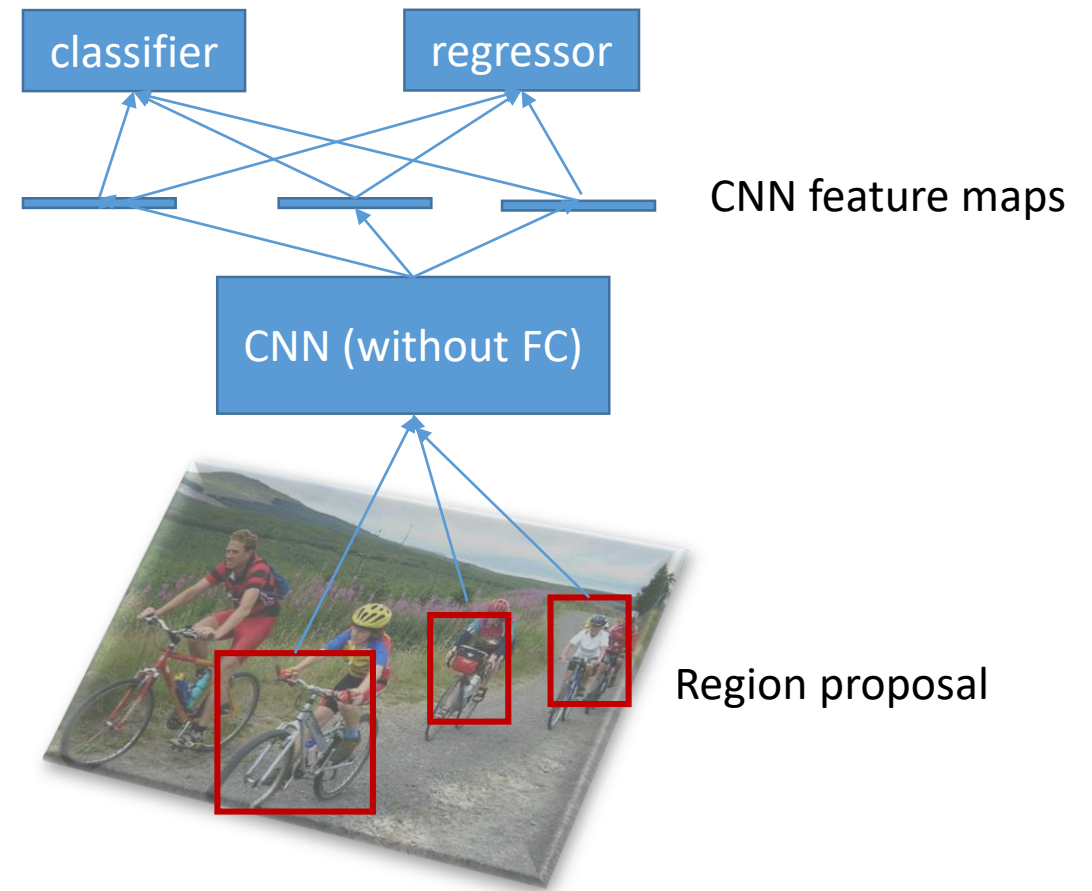
# R-CNN

- Generate region proposal
  - Candidate object regions
  - Using existing methods
- Extract CNN feature
  - For each region
  - Forward-propagate each region via CNN
- Predict label for each region
  - Like image classification
  - Linear layer + softmax



# R-CNN

- Generate region proposal
  - Candidate object regions
  - Using existing methods
- Extract CNN feature
  - For each region
  - Forward-propagate each region via CNN
- Predict label for each region
  - Like image classification
  - Linear layer + softmax
- Regress bounding box for each region
  - Linear regression for each value
  - 4 values (coordinates, or  $x, y, h, w$ )



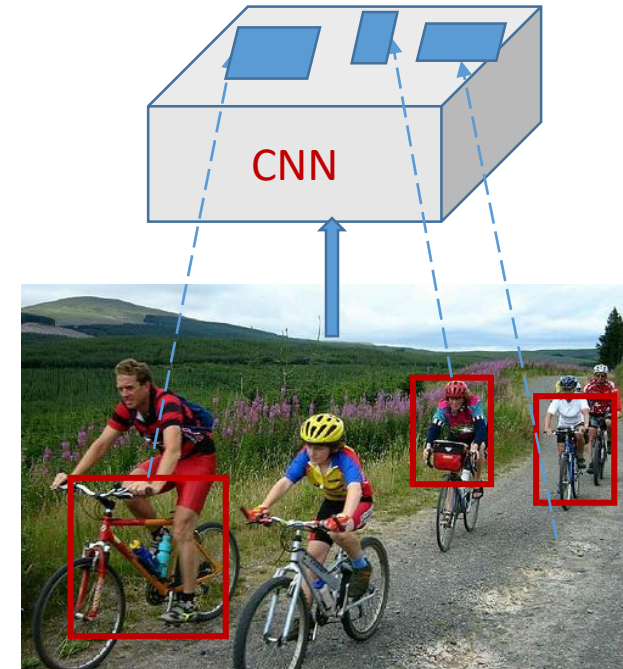


# R-CNN

- Slow
  - Too many region proposals~2000
  - Each has to go through the CNN
- Training is ad-hoc
  - Fine tune the CNN for the target dataset for image classification
  - Train label classifier for regions
  - Train SVM regressor for bounding box

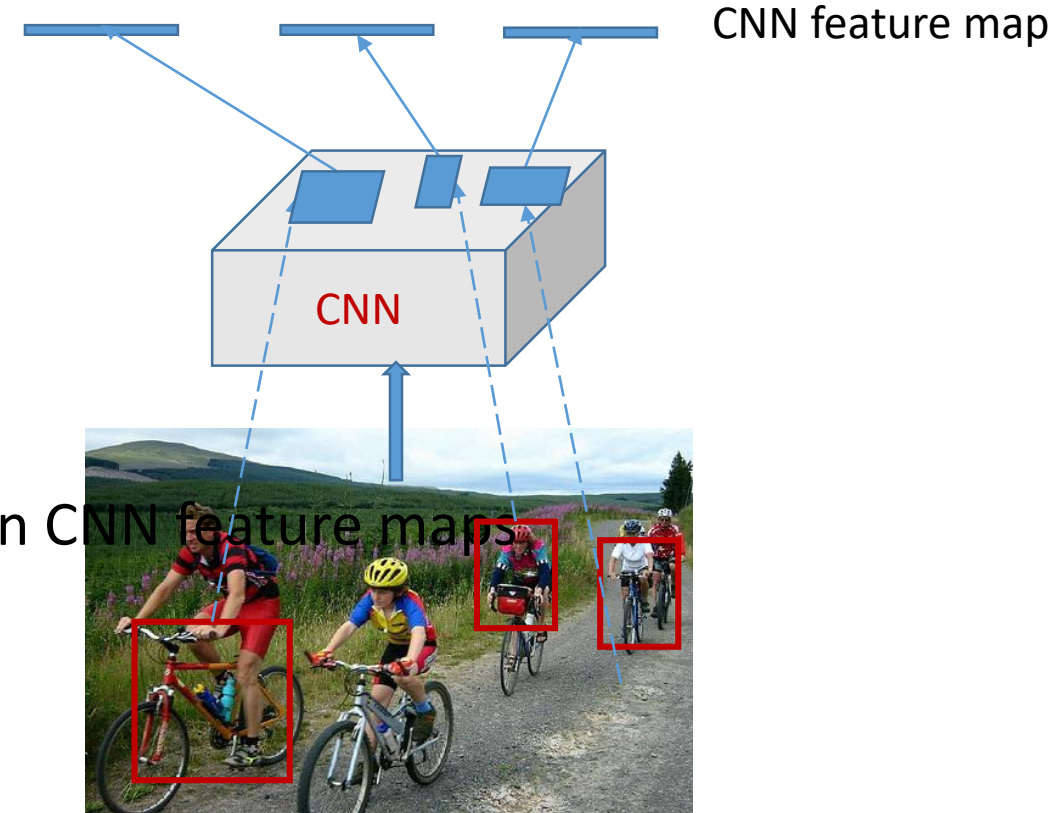
# Fast R-CNN[2]

- Generate region proposal
  - Candidate object regions
  - Using existing methods
- Extract CNN feature
  - For the **whole input image**
  - Forward-propagate the whole image via CNN



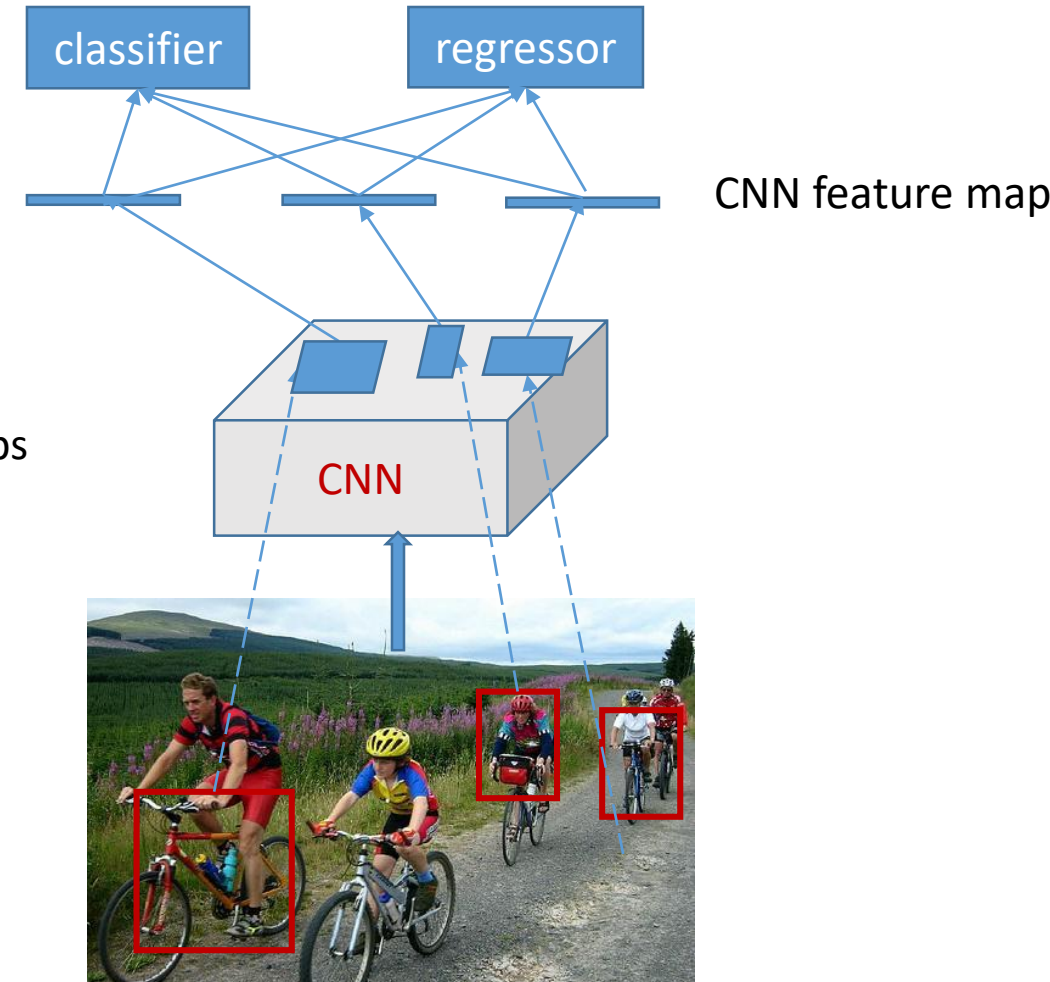
# Fast R-CNN

- Generate region proposal
  - Candidate object regions
  - Using existing methods
- Extract CNN feature
  - For the **whole input image**
  - Forward-propagate the whole image via CNN
- Get CNN feature for each region
  - Using the region coordinates to locate areas in CNN feature maps
  - Extract the CNN feature from the areas



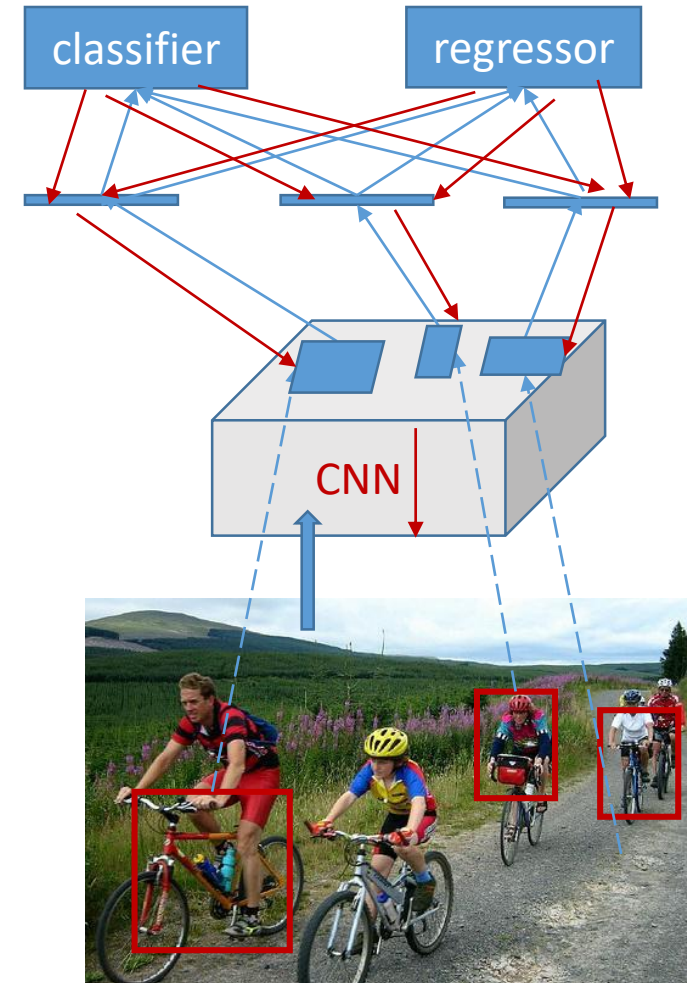
# Fast R-CNN

- Generate region proposal
  - Candidate object regions
  - Using existing methods
- Extract CNN feature
  - For the **whole input image**
  - Forward-propagate the whole image via CNN
- Get CNN feature for each region
  - Using the region coordinates to locate areas in CNN feature maps
  - Extract the CNN feature from the areas
- Predict label for each region
  - Like image classification
  - Linear layer + softmax
- Regress bounding box for each region
  - Linear regression for each value
  - 4 values (coordinates, or  $x, y, h, w$ )



# Fast RCNN

- Run CNN forwarding once
- End-to-end training
  - Softmax classifier
  - Linear regressor for bounding box

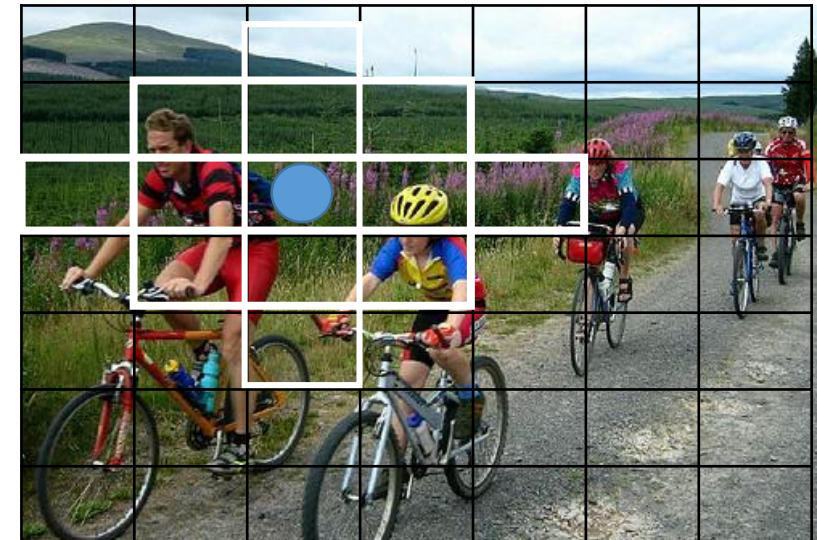




# YOLO [3]

- Extract CNN feature for the whole image
- Divide the input image into fixed number of grids (e.g. 7x7)
  - For each cell
    - generate B (e.g. 3) bounding candidate boxes with different aspect ratios
    - Get the feature of each candidate box
    - Use Softmx to do label classification; linear regressor for coordinates refinement.
- NO Offline candidate region generation
- [Notebook](#)

Fast!  
Not very accurate!



# Image segmentation

- Label each pixel with a class
- Training label
  - A class (index) per pixel
- Prediction output
  - For each pixel, a probability vector (one per class)



Training label:

(0, 0) bicycle ...

(200, 80) people (200, 81) people

Prediction output:

(0, 0) background 0.9; mountain: 0.1

...

(200, 80) people 0.8; bicycle 0.1; tree: 0.1

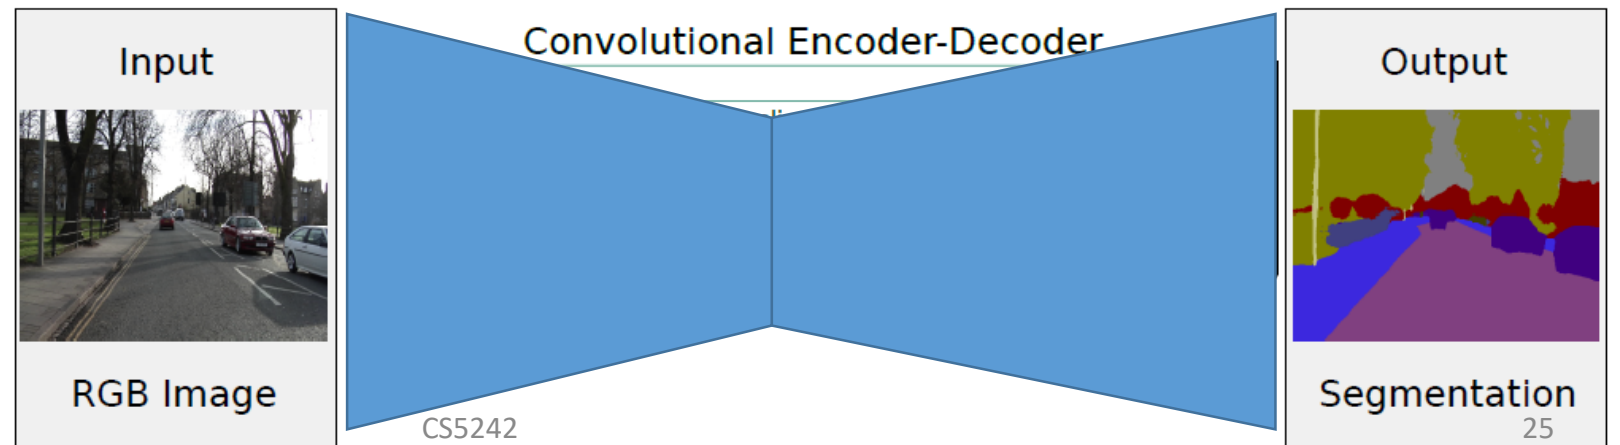
# Image segmentation

- Applications
  - Medical image analysis
  - [Self-driving car](#)
- Evaluation [1]
  - Matched pixels = the predicted class of a pixel is the truth class
  - Mean IoU = average over all classes  $\{\text{\#matched pixels} / (\text{truth pixels} \cup \text{predicted pixels})\}$



# Image segmentation

- Solution
  - Encoder to extract a semantic-rich representation
    - For label prediction
    - Subsampling by (pooling or convolution with stride  $> 1$ )
  - Decoder to incorporate location information
    - To generate a final feature map as the same size as the input
    - Upsampling
- Loss
  - Softmax loss for each pixel



# Image segmentation

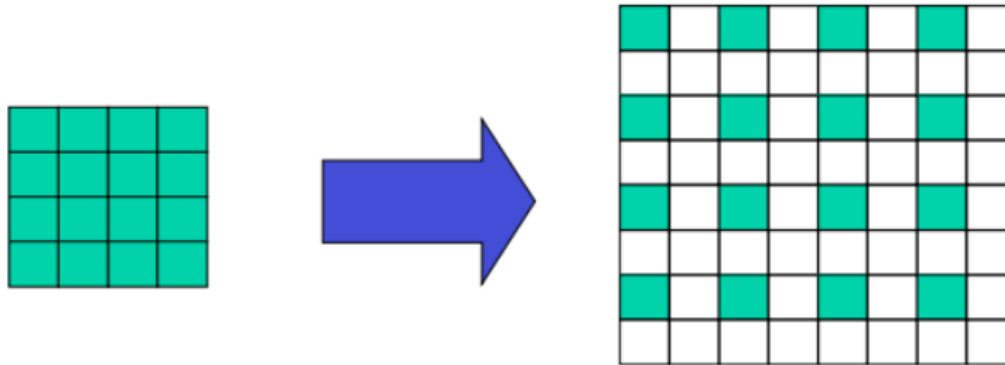
- Upsampling
  - Nearest neighbour

3	2
0	1

3	3	2	2
3	3	2	2
0	0	1	1
0	0	1	1

# Image segmentation

- Bilinear upsampling



- The empty pixels are initially set to 0
- Convolve with a (Gaussian, or another) filter
- If the filter sums to 1, multiply the result by 4
  - $\frac{3}{4}$  of the new image was initially 0

Original image:  x 10



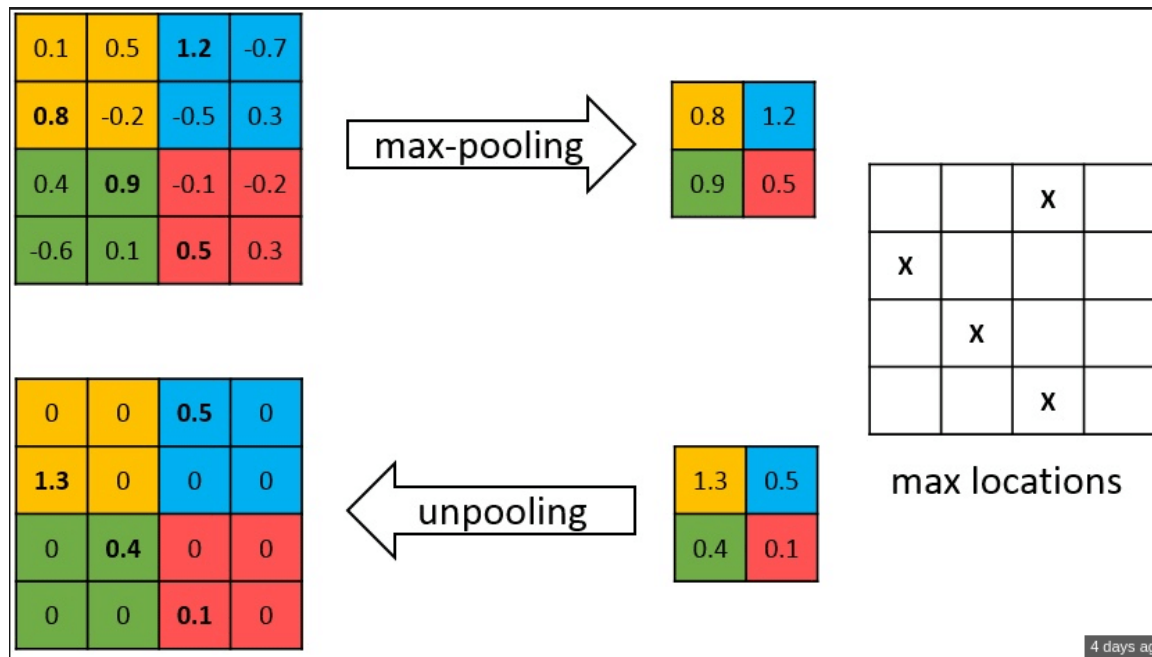
Nearest-neighbor interpolation



Bilinear interpolation

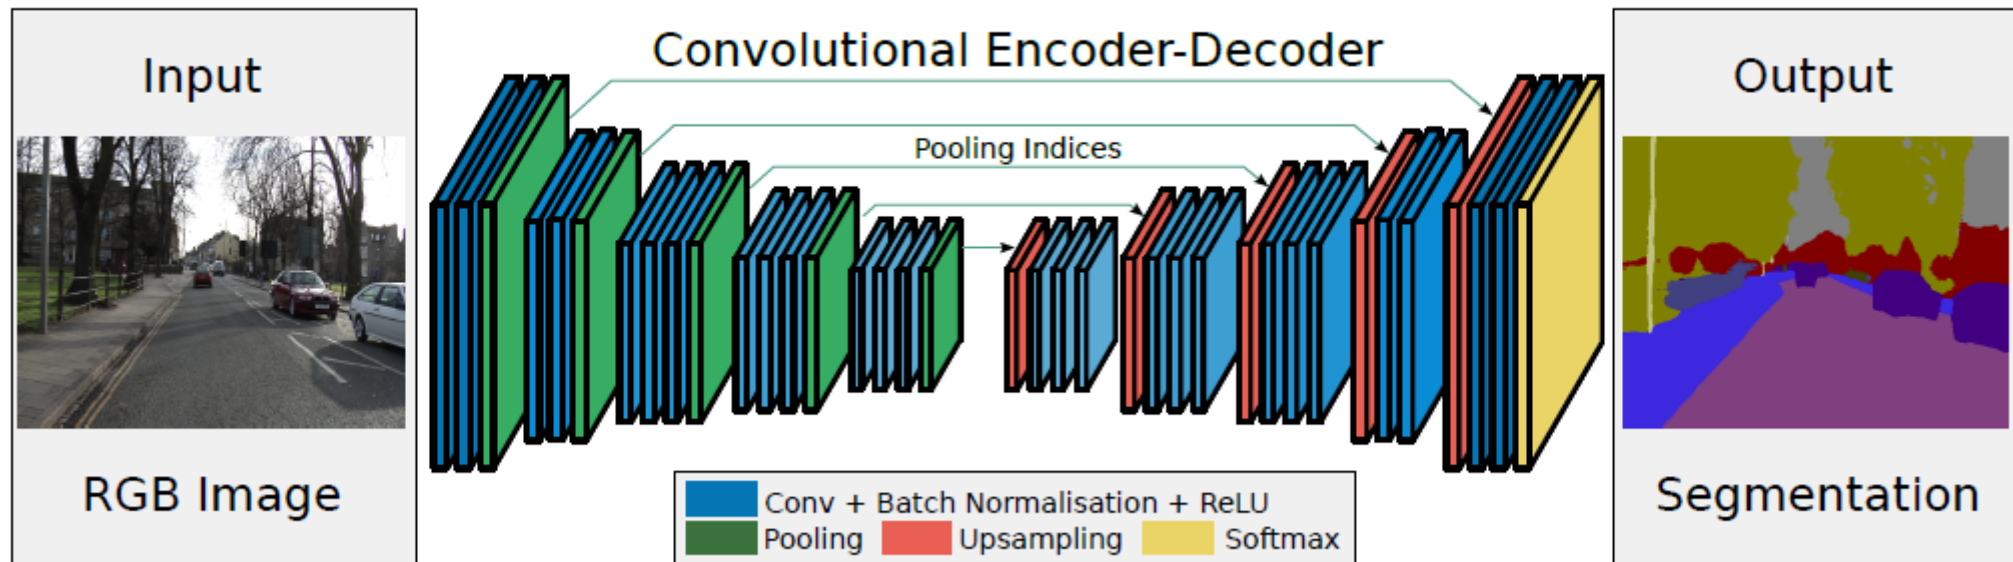
# Image segmentation

- Max unpooling



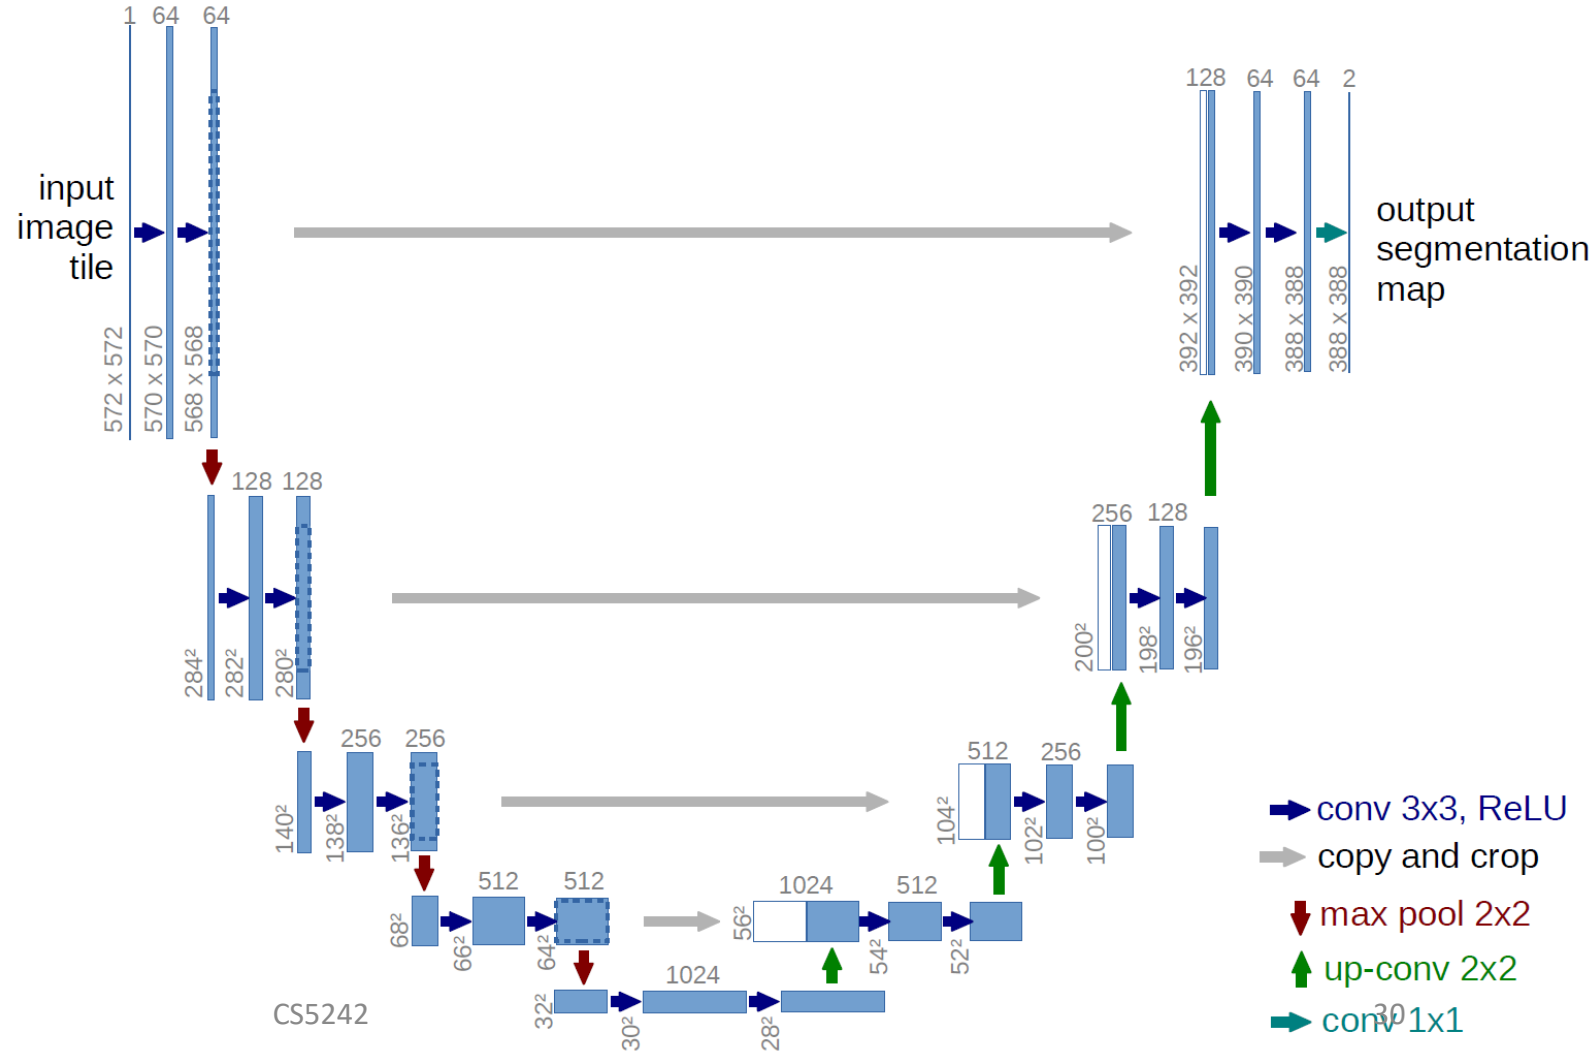
# Image segmentation

- SegNet [4]
  - Max unpooling



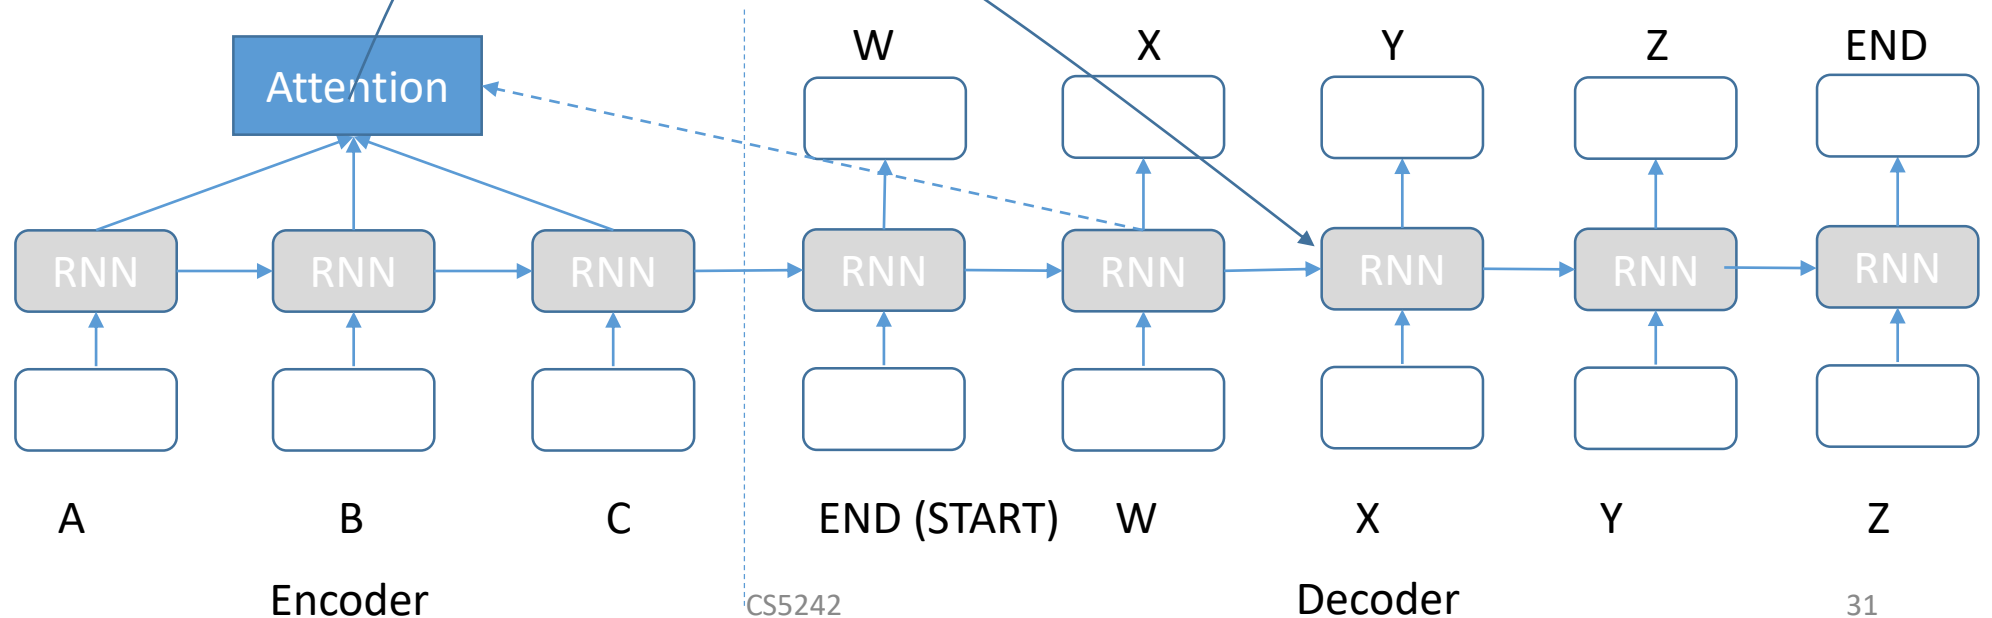
# U-Net[5]

- Prototxt visualization
- Input and output
  - Different size
  - Due to valid padding
- Examples
- 1, 2



# Attention modelling [6]

- Each output word depends on
  - All input words (hidden states), each with different contribution
    - Singapore MRT is not stable
    - 新加坡 地铁 不 稳定



# Attention modelling [6]

- Encoder
  - Input  $a=[0.1,0,1]$ ,  $b=[1,0.1,0]$ ,  $c=[0.2,0.3,1]$
  - Hidden representation (vector)
    - $h_1, h_2, h_3$
    - $h_1=\tanh(Ua+Wh_0)$
    - $h_2=\tanh(Ub+Wh_1)$
    - $h_3=\tanh(Uc+Wh_2)$



- Decoder
  - Hidden state  $s_0 = [0,0,0]$  or  $h_3$ 
    - To compute the weights of  $h_1, h_2, h_3$  for computing  $s_1$ 
      - $e_{11} = a(s_0, h_1), e_{12}=a(s_0, h_2), e_{13}=a(s_0, h_3)$
      - $a(s_0, h_1) = v^T \tanh(W_a s_0 + U_a h_1)$
      - $a(s_0, h_2) = v^T \tanh(W_a s_0 + U_a h_2)$
      - $a(s_0, h_3) = v^T \tanh(W_a s_0 + U_a h_3)$
      - $k_{11} = \exp(e_{11}) / (\exp(e_{11}) + \exp(e_{12}) + \exp(e_{13}))$
      - $k_{12} = \exp(e_{12}) / (\exp(e_{11}) + \exp(e_{12}) + \exp(e_{13}))$
      - $k_{13} = \exp(e_{13}) / (\exp(e_{11}) + \exp(e_{12}) + \exp(e_{13}))$
      - $c_{11} = k_{11}h_1 + k_{12}h_2 + k_{13}h_3$
  - $s_t = (1 - z_t) \circ s_{t-1} + z_t \circ \tilde{s}_t$
  - $\tilde{s}_t = \tanh(W([r_t \circ s_{t-1}, Ey_{t-1}, c_t]))$
  - $r_t = \sigma(W_r[s_{t-1}, Ey_{t-1}, c_t])$
  - $z_t = \sigma(W_z[s_{t-1}, Ey_{t-1}, c_t])$

# More Saturday sessions

- Other topics?
- Assignments answers will be uploaded

# Reference

- [1] Yunchao Gong, Yangqing Jia, Thomas Leung, Alexander Toshev, Sergey Ioffe. Deep convolutional ranking for multilabel image annotation. <https://arxiv.org/pdf/1312.4894.pdf>
- [2] Ross Girshick. Fast R-CNN. <https://arxiv.org/abs/1504.08083>
- [3] Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi. You Only Look Once: Unified, Real-Time Object Detection. <https://arxiv.org/abs/1506.02640>
- [4] SegNet: A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation. Vijay Badrinarayanan, Alex Kendall, Roberto Cipolla. 2016
- [5] U-Net: Convolutional Networks for Biomedical Image Segmentation. Olaf Ronneberger, Philipp Fischer, Thomas Brox. 2015
- [6] <https://machinelearningmastery.com/how-does-attention-work-in-encoder-decoder-recurrent-neural-networks/>
- [7] R. Girshick, J. Donahue, T. Darrell, and J. Malik, “Rich feature hierarchies for accurate object detection and semantic segmentation,” in IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2014