



Recurrent Neural Networks (RNN)

CS5242

Lee Hwee Kuan & **Wang Wei**

Teaching assistant:

Connie Kou Khor Li, Ji Xin, Ouyang Kun

cs5242@comp.nus.edu.sg

Mid-term survey

- Topics
 - Basics (output size, computation cost) VS advanced stuff (architectures)
 - More software/implementation practices VS more mathematics
 - More Introduction/background/motivation
- Slides
 - More explanation/notes
 - Fonts, figures
- Workload
 - Assignment and projects
- Voice & Pronunciation

Announcement

- Quiz 20%
 - Oct. 26, 18:30-19:30, open book
- NO assignment 4
 - Assignment 1: 10%
 - Assignment 2: 15%, Due date: 22 Oct. 11:59PM (extended)
 - Assignment 3: 15%
- Saturday session
 - 15:30-17:30
 - Lecture room (I3 auditorium room)
 - IVLE survey

Roadmap



Intended learning outcomes

01

Understand the properties of RNN compared with feed-forward NN

02

Implement the BP of vanilla RNN


03

Know the problem of vanilla RNN and the properties of LSTM/GRU

04

Train RNN (vanilla/LSTM/GRU) for language modelling

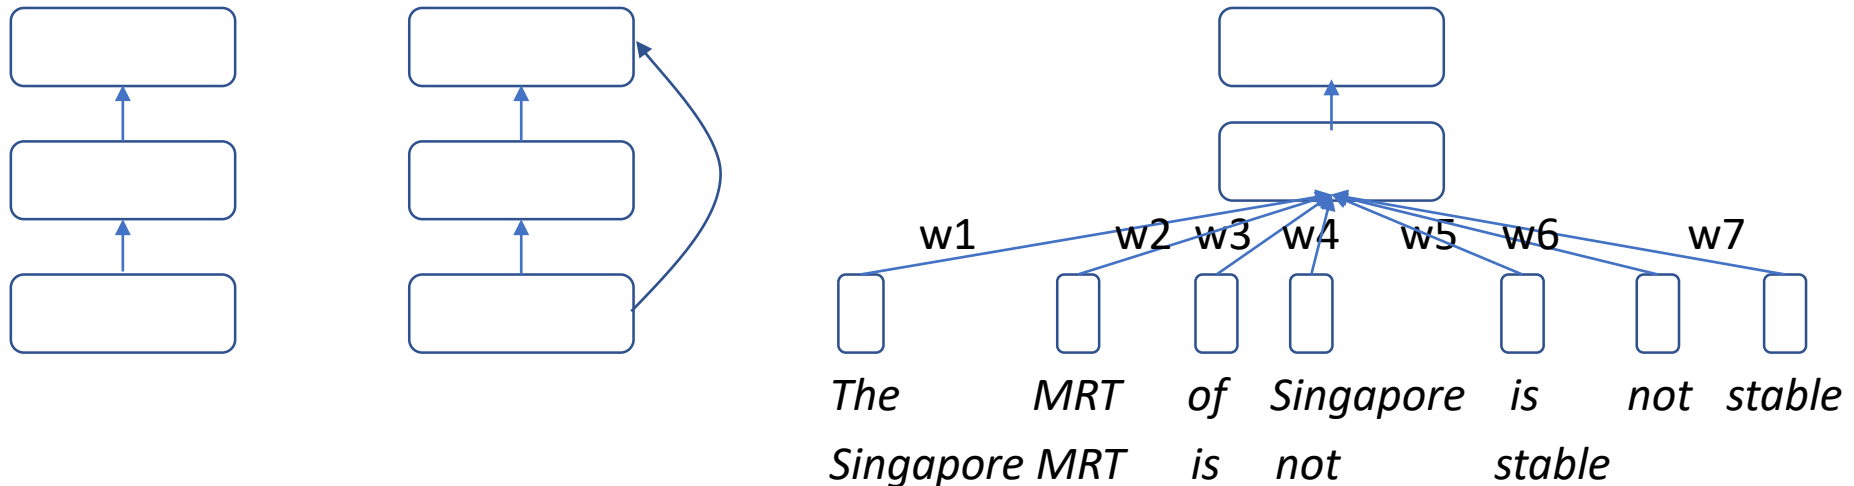
Motivation



From feed-forward NN to RNN

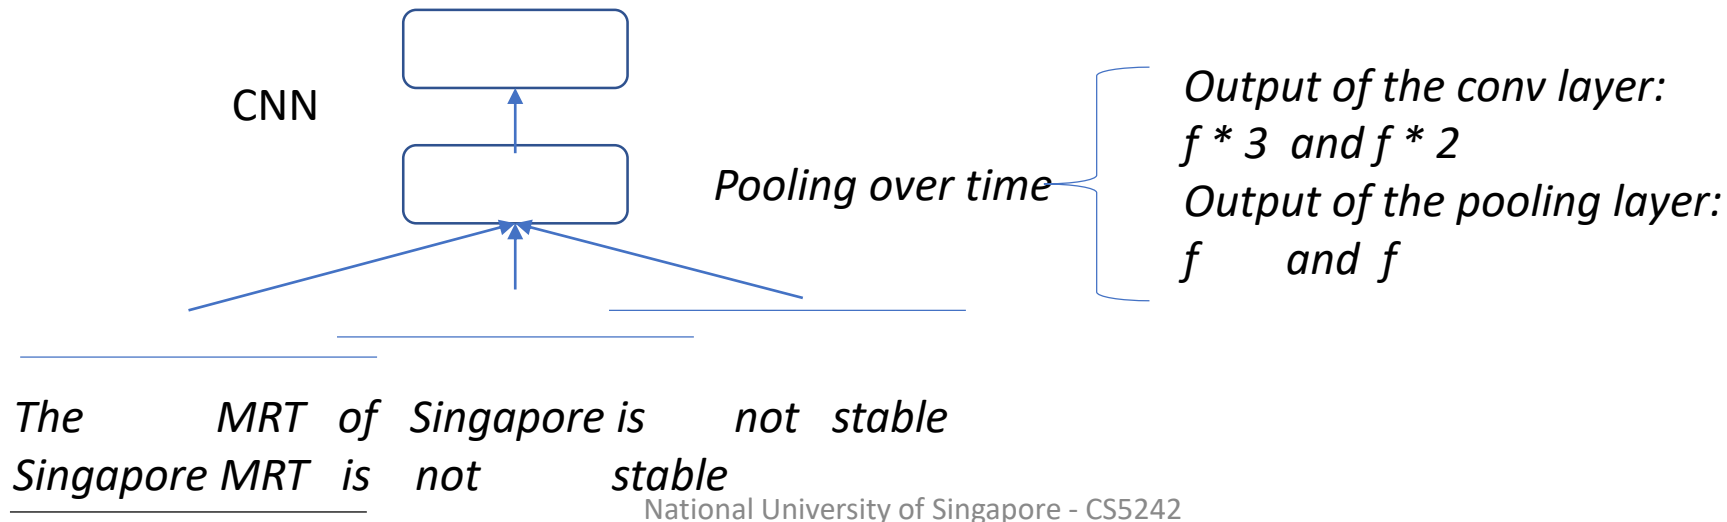
- Feed-forward NN (acyclic)
 - Accept single/static input sample, e.g. image
 - Not good at processing a sequence of data
 - E.g. a sentence of words for sentiment analysis; **how to do it using MLP?**

Feed-forward NN



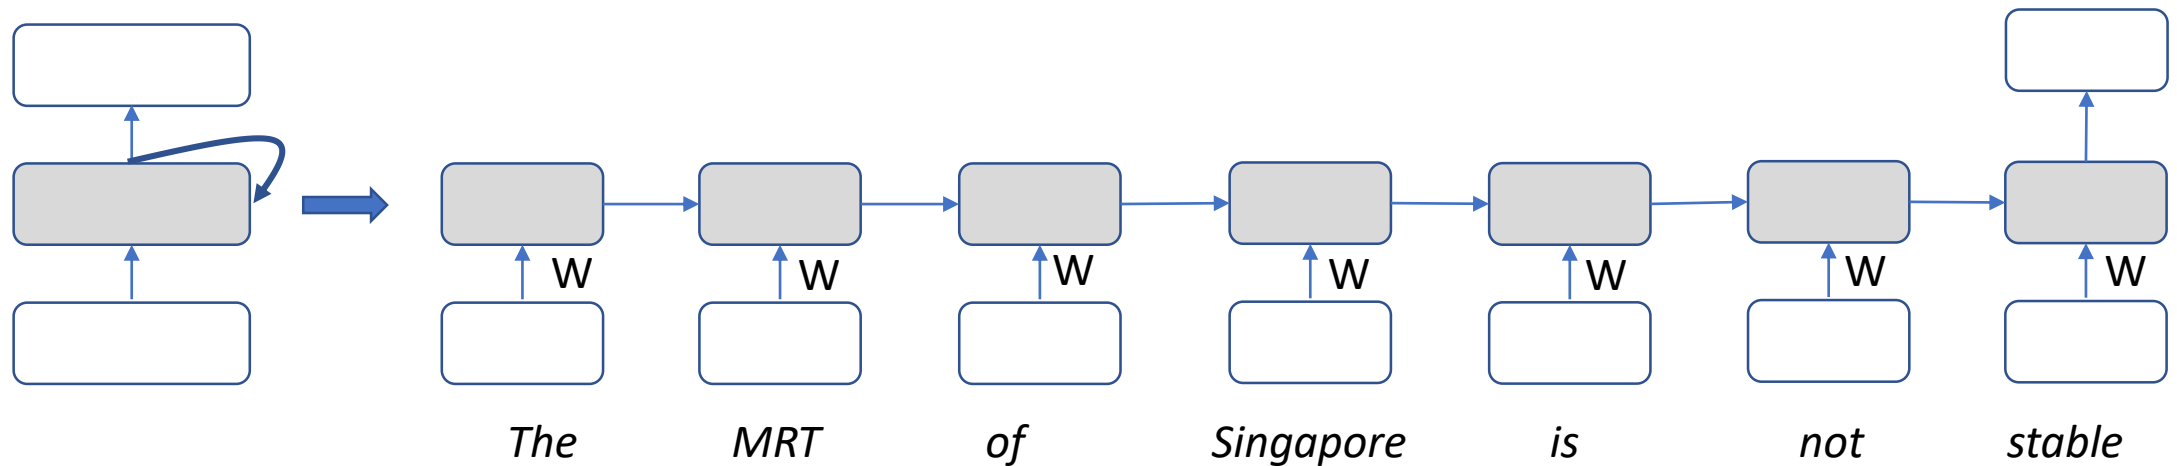
From feed-forward NN to RNN

- Feed-forward NN (acyclic)
 - Accept single/static input sample, e.g. image
 - Not good at processing a sequence of data
 - CNN's receptive fields share the parameters (i.e. kernels)
 - Kernel size typically > 1 \rightarrow words within the receptive field are processed differently
 - "MRT of Singapore" \neq "Singapore MRT is"



From feed-forward NN to RNN


- Feed-forward NN (acyclic)
- RNN
 - Accept dynamic/sequence data (length not fixed)
 - Words are processed in the same way recurrently
 - # unfold units = input sequence length
 - weights are tied



RNN

- Applications
 - Language modelling
 - Predict the next word given the previous words in a sentence
 - Machine translation
 - Translate the input English sentence to French
 - Speech recognition
 - Recognize and translate spoken language into text
 - Question answering
 - Generate text answers for (simple) questions
 - Etc.

Vanilla RNN

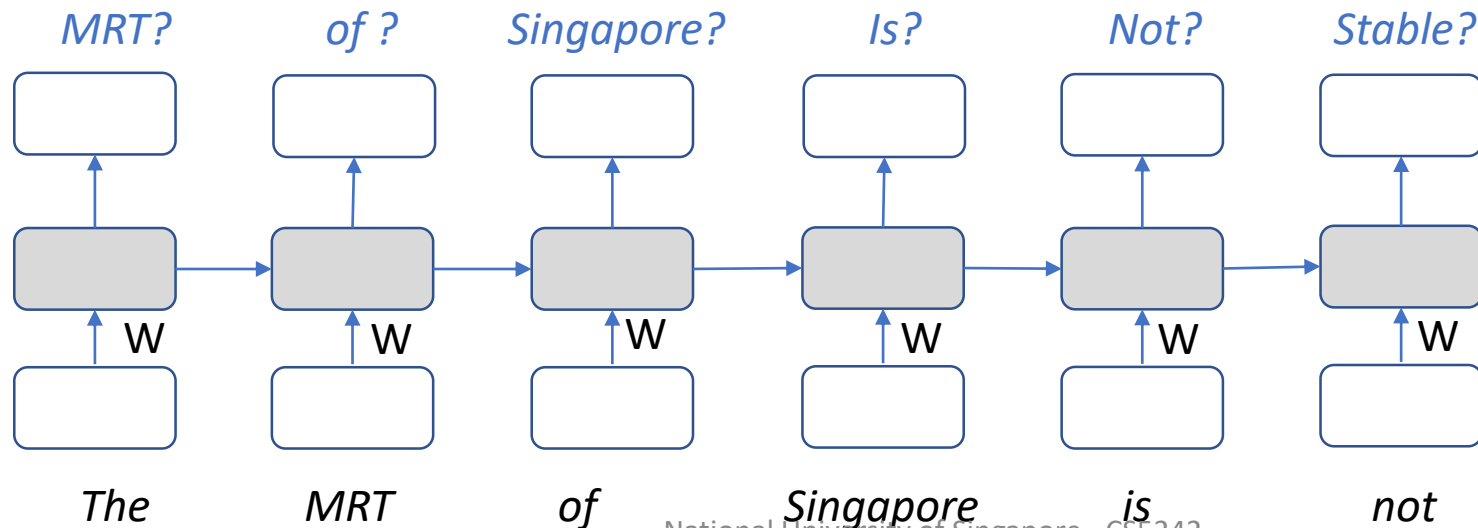


Language modelling example

- Given a corpus of text (e.g. sentences), model the probability of a sentence (i.e. a sequence of words)
 - $P(x_1, x_2, \dots, x_n)$
 - Useful for many applications involving text/sentence generation?
 - Machine translation, speech recognition, question answering, etc.
 - $P(\text{"Singapore MRT is not stable"}) > P(\text{"Singapore MRT is not NUS"})$
 - $P(\text{"Singapore MRT is not stable"}) > P(\text{"Singapore is MRT not stable"})$
 - Refer to [5] for traditional approaches for this problem


Language modelling example

- $P(x_1, x_2, \dots, x_n) = \prod_t P(x_t | x_{t-1}, x_{t-2}, \dots, x_1)$
 - $\log P(x_1, x_2, \dots, x_n) = \sum_t \log P(x_t | x_{t-1}, x_{t-2}, \dots, x_1) \rightarrow ?$
 - Maximize log-likelihood == minimize cross-entropy
 - Train classifiers to predict the next word given the preceding words
 - $P(\text{MRT} | \text{The})$, $P(\text{of} | \text{The MRT})$, $P(\text{Singapore} | \text{The MRT of})$, ...




Hidden Layer

- Denote the vocabulary of all words as V
 - {MRT:0, Singapore:1, Stable:2, ...}.
- Represent the word at position t as a column vector $x_t \in R^d$,
 - E.g., word vectors [2,3]; d is word vector length, e.g. 32, 64, 128.
 - Retrieve the word vector from the downloaded file using the word as key




Hidden Layer

- Denote the vocabulary of all words as V
 - $\{\text{MRT:0, Singapore:1, Stable:2, ...}\}$.
- Represent the word at position t as a column vector $x_t \in R^d$,
 - E.g., word vectors $[2,3]$; d is word vector length, e.g. 32, 64, 128.
- Denote the hidden layer at position t as $h_t \in R^k$
 - k is defined by users
 - $h_t = f(h_{t-1}, x_t | \theta)$
 - $a_t = Ux_t + Wh_{t-1} + b$,
 - $\theta = \{U \in R^{k \times d}, W \in R^{k \times k}, b \in R^k\}$
 - $h_t = \tanh(a_t), h_t \in R^k$



Output

- $o_t = Vh_t + c$,
 - $V \in R^{|V| \times k}, c \in R^{|V|}, o_t \in R^{|V|}$
- $y_t = \text{softmax}(o_t)$
 - If $|V|$ is very large, the prediction layer needs special optimization [6]
 - $y_t \in R^{|V|}$, a probability vector

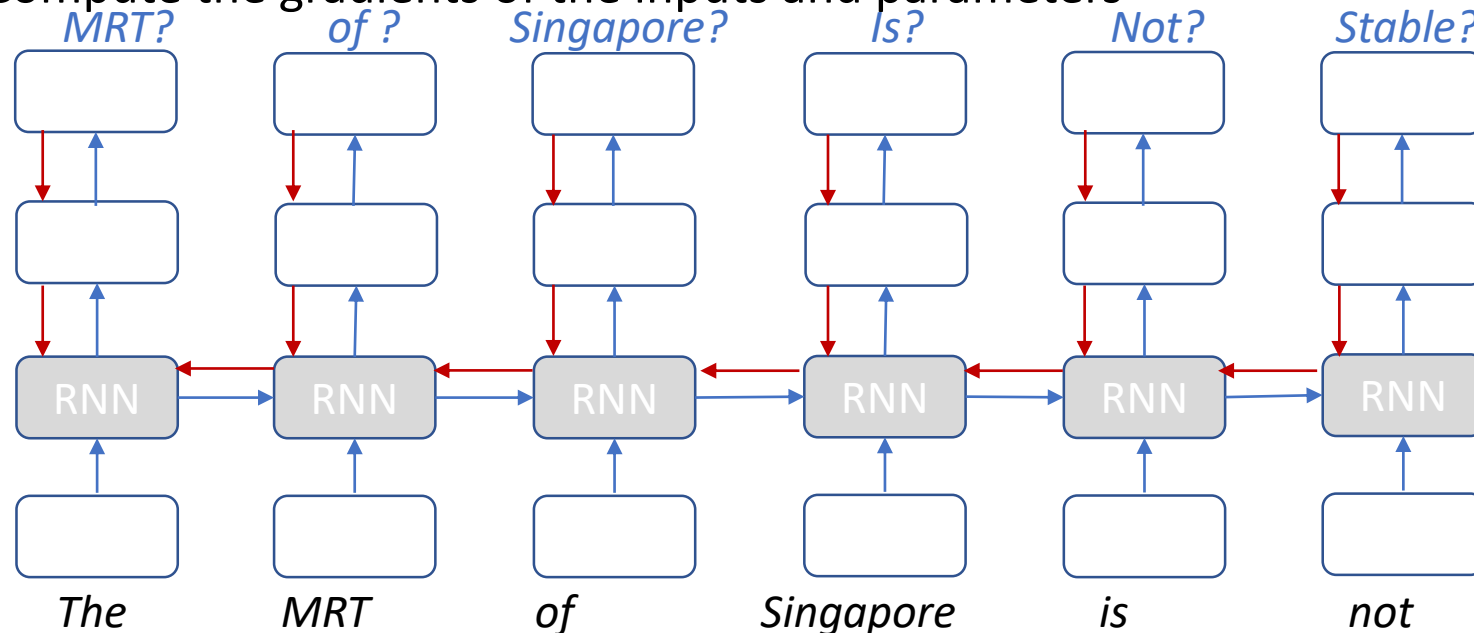


Training

- Given a corpus of text data (e.g. sentences), train the parameters of the RNN.
- Objective:
 - Maximize $\log P(x_1, x_2, \dots, x_n) = \sum_t \log P(x_t | x_{t-1}, x_{t-2}, \dots, x_1)$
 - \rightarrow minimize the cross-entropy loss at each position t , denoted as L_t
 - $L = \sum_t L_t$
- SGD
 - For each data sample (e.g. a sentence)
 - Compute the gradients of each parameter
 - Update the parameters by $\theta = \theta - \alpha \times \frac{\partial L}{\partial \theta}$

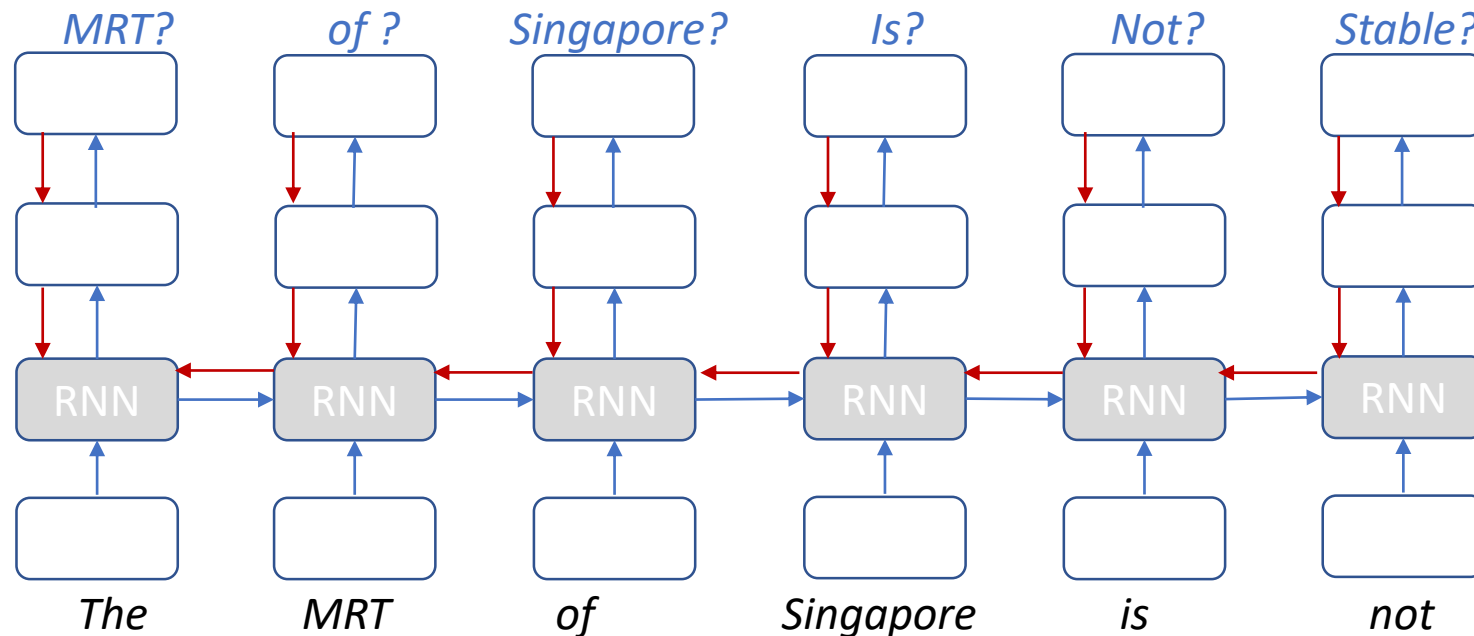
Back-propagation through time (BPTT)

- Back-propagation for each position as normal
 - From the cross-entropy to the RNN layer
 - For each RNN layer
 - aggregate the gradients from the top layer and the right layer
 - Compute the gradients of the inputs and parameters



Back-propagation through time (BPTT)

- Back-propagation for each position as normal
 - From the cross-entropy to the RNN layer
 - For each RNN layer
 - For each parameter, aggregate the gradient across all positions



Back-propagation through time (BPTT)

- Forward


- $a_t = Ux_t + Wh_{t-1} + b,$
- $h_t = \tanh(a_t)$
- $o_t = Vh_t + c, y_t = \text{softmax}(o_t)$

- Softmax+cross-entropy

- $\frac{\partial L_t}{\partial o_t} = y_t - l_t, l_t \in \{0,1\}^{|V|}$, the ground truth vector

- Dense

- $\frac{\partial L_t}{\partial h_t} = V^T \frac{\partial L_t}{\partial o_t}$
- $\frac{\partial L_t}{\partial V_t} = \left(\frac{\partial L_t}{\partial o_t} \right) (h_t)^T, \frac{\partial L_t}{\partial c_t} = \frac{\partial L_t}{\partial o_t}$, gradients of V and c from t-th position



Back-propagation through time (BPTT)

- Forward

- $a_t = Ux_t + Wh_{t-1} + b,$
- $h_t = \tanh(a_t)$
- $o_t = Vh_t + c$

- RNN layer

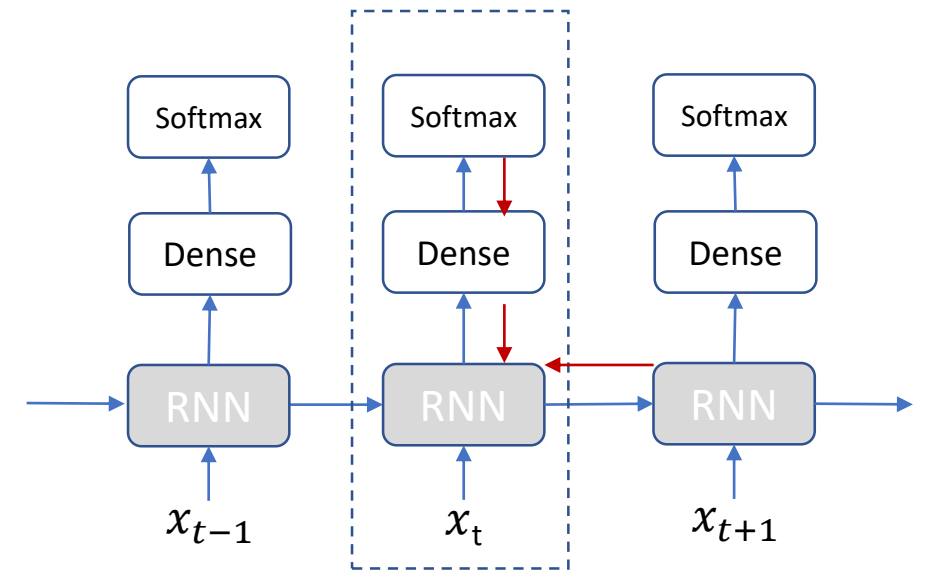
$$\frac{\partial L}{\partial h_t} = \frac{\partial L_t}{\partial h_t} + \boxed{\frac{\partial L_{t+1}}{\partial h_t} + \dots + \frac{\partial L_n}{\partial h_t}} = \frac{\partial L_t}{\partial h_t} + \frac{\partial L_{t+}}{\partial h_t}$$

$$\frac{\partial L}{\partial a_t} = \frac{\partial L}{\partial h_t} \times (1 - h_t^2)$$

$$\frac{\partial L}{\partial U_t} = \frac{\partial L}{\partial a_t} x_t^T, \frac{\partial L}{\partial W_t} = \frac{\partial L}{\partial a_t} h_{t-1}^T, \frac{\partial L}{\partial b_t} = \frac{\partial L}{\partial a_t}, \text{gradients of } U, W, b \text{ from position } t$$

$$\frac{\partial L_{(t-1)+}}{\partial h_{t-1}} = W^T \frac{\partial L}{\partial a_t}$$

$$\frac{\partial L}{\partial \theta} = \sum_t \frac{\partial L}{\partial \theta_t}$$



Back-propagation through time (BPTT)

- Gradient vanishing/exploding

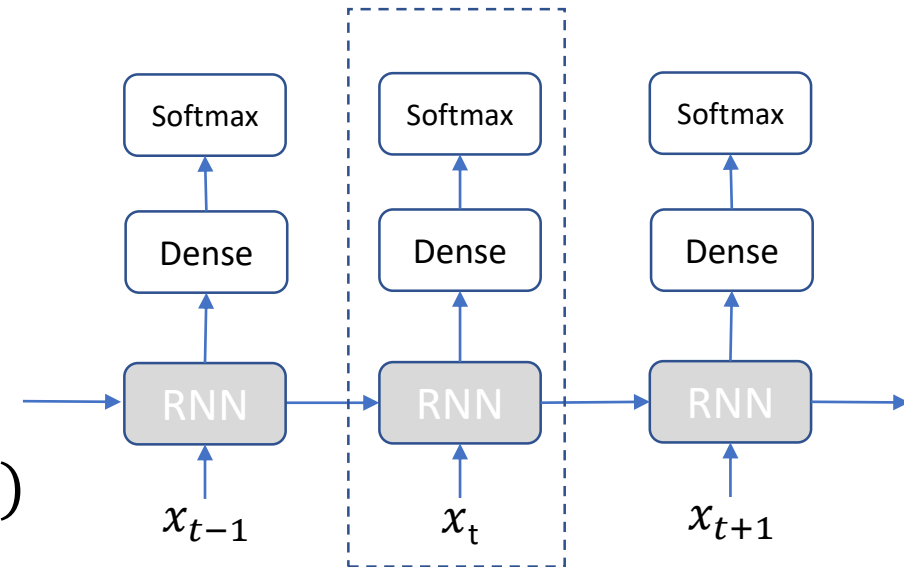
- $$\frac{\partial L}{\partial h_t} = \frac{\partial L_t}{\partial h_t} + \frac{\partial L_{t+1}}{\partial h_t} + \dots + \frac{\partial L_n}{\partial h_t} = \frac{\partial L_t}{\partial h_t} + \frac{\partial L_{t+}}{\partial h_t}$$

- $$\frac{\partial L}{\partial a_t} = \frac{\partial L}{\partial h_t} \times (1 - h_t^2)$$

- $$\rightarrow \frac{\partial L_{(t-1)+}}{\partial h_{t-1}} = W^T \frac{\partial L}{\partial a_t} = W^T \frac{\partial L}{\partial h_t} \times (1 - h_t^2)$$

- $$= W^T \left(\frac{\partial L_t}{\partial h_t} + \frac{\partial L_{t+}}{\partial h_t} \right) \times (1 - h_t^2)$$


- $$\rightarrow \frac{\partial L_{(t-1)+}}{\partial h_{t-1}} \leftarrow W^T \frac{\partial L_{t+}}{\partial h_t} \dots \leftarrow (W^T)^k \frac{\partial L_{(t+k)+}}{\partial h_t}$$



Gradients from right most positions vanish when back-propagated to the left-most positions

Back-propagation through time (BPTT)

- $\frac{\partial L_{(t-1)+}}{\partial h_{t-1}} \leftarrow W^T \frac{\partial L_{t+}}{\partial h_t} \dots \leftarrow (W^T)^k \frac{\partial L_{(t+k)+}}{\partial h_t}$
 - If $|W|$ is small, gradient vanishing
 - The losses after position $t+k$ have little influence for the RNN layer at $t-1$ if k is large
 - Cannot capture long-term relationship
 - “The **red line** went down last night, which is why there are many tweets about ___ “(red line).
 - If $|W|$ is large, gradient exploding
- Solutions
 - Gradient vanishing?
 - Careful initialization
 - Identity matrix with ReLU as the activation function[7]
 - Skip-connections
 - leaky units \rightarrow LSTM and GRU
 - $h_t = \gamma h_{t-1} + (1 - \gamma) \tanh(Ux_t + Wh_{t-1} + b)$
 - Gradient exploding?
 - Gradient clipping



Gradient Clipping

- $W = W - \alpha \times \frac{\partial L}{\partial W}$
- Hard clipping
 - For each value of $\frac{\partial L}{\partial W}$, if it is larger than a threshold μ , set it to be μ
- Normalization (L2)
 - $g = \frac{\partial L}{\partial W}$
 - If $|g| > \mu$, $g = \frac{\mu}{|g|} g$

Mini-batch SGD

- SGD uses a single sample per iteration
- Mini-batch SGD uses multiple samples per iteration
 - To accelerate the processing by matrix (batch) operations
 - Different sentences have different lengths, e.g.

Singapore MRT is not stable
Chicken rice is very popular in Singapore
It is hot


Word to index



0 2 1 5 6
7 8 1 4 8 9 0
10 1 11

Mini-batch SGD

- Solution?
 - Truncate the sentences into the same fixed length



Mini-batch SGD

- Solution

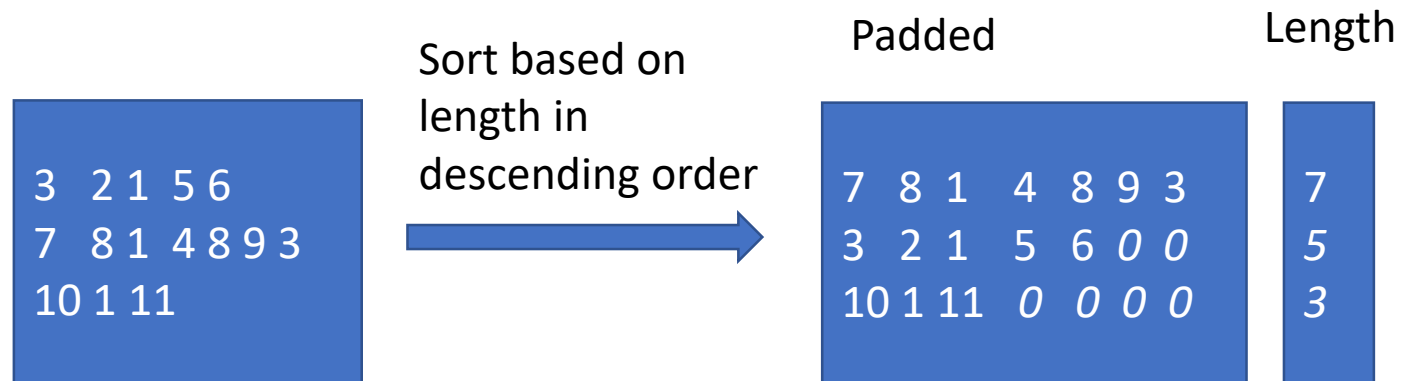
- Truncate the sentences into the same fixed length

- Padding

- E.g. [PyTorch](#)


`pack = torch.nn.utils.rnn.pack_padded_sequence(batch_in, seq_lengths, batch_first=True)`

- Index 0 is for a special 'PAD' symbol. Index of words in the vocabulary starts from 1.




Other tricks for training

- Adaptive learning rate
 - E.g. Adam, RMSProp
- Normalizing the losses
 - $L = \sum_t L_t \rightarrow L = \frac{1}{n \sum_t L_t}$
- Use gated RNN units
 - LSTM or GRU (not introduced yet)
- Stack multiple RNN layers
 - As shown by the right figure
- Layer normalization [8, 10]
 - Applied before activation function
- Recurrent Dropout [9, 10]



Other tricks for training

- Learn the initial state h_0 [11]
 - Typically, we set h_0 to be a all 0 vector
 - It can also be learned like a bias vector
 - computing the gradient and then apply SGD update
- Trunked BPTT
 - Some sentences are very long, e.g. > 1000 positions.
 - Split the sentence into shorter sub-sentences, e.g. 200
 - Each sub-sentence is a new training sample
 - Use the last hidden vector (h_n) of the previous sub-sentence as the initial state h_0 for the next sub-sentence



Reference

- [1] <https://www.quora.com/What-are-differences-between-recurrent-neural-network-language-model-hidden-markov-model-and-n-gram-language-model>
- [2] <https://code.google.com/archive/p/word2vec/>
- [3] <https://nlp.stanford.edu/projects/glove/>
- [4] Klaus Greff, Rupesh Kumar Srivastava, Jan Koutník, Bas R. Steunebrink, Jürgen Schmidhuber. LSTM: A Search Space Odyssey. <https://arxiv.org/abs/1503.04069>
- [5] <http://web.stanford.edu/class/cs224n/lectures/cs224n-2017-lecture8.pdf>
- [6] <http://www.deeplearningbook.org/contents/applications.html> (12.4.3)
- [7] Quoc V. Le, Navdeep Jaitly, Geoffrey E. Hinton. A Simple Way to Initialize Recurrent Networks of Rectified Linear Units. 2015. arxiv.org/abs/1504.00941v2
- [8] "Layer Normalization" Jimmy Lei Ba, Jamie Ryan Kiros, Geoffrey E. Hinton. <https://arxiv.org/abs/1607.06450>.
- [9] "Recurrent Dropout without Memory Loss" Stanislaw Semeniuta, Aliaksei Severyn, Erhardt Barth. <https://arxiv.org/abs/1603.05118>
- [10] https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/LayerNormBasicLSTMCell
- [11] <https://r2rt.com/non-zero-initial-states-for-recurrent-neural-networks.html>
- [12] LSTM: A Search Space Odyssey. Klaus Greff, Rupesh Kumar Srivastava, Jan Koutník, Bas R. Steunebrink, Jürgen Schmidhuber. <https://arxiv.org/abs/1503.04069>
- [13] <https://github.com/karpathy/char-rnn/issues/138#issuecomment-162763435>
- <https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>
- <https://danijar.com/tips-for-training-recurrent-neural-networks/>